

SIDDHANT GUPTA

DATA ENGINEER

(+91) 8851728513 | siddhantg2002@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Bengaluru, India | Open to Remote / Relocation

SUMMARY

Data Engineer with 1.5+ years building HIPAA-compliant healthcare data pipelines at scale on Databricks and AWS. Specialises in Medallion Architecture, PySpark, and real-time Kafka streaming. Experienced in LLM-assisted SQL migration, config-driven ETL frameworks, and CI/CD-integrated test automation. Shipped open-source developer tools recognised in Top 50 on Peerlist.

SKILLS

Languages: Python, SQL, JavaScript/TypeScript, Bash

Data & Cloud: PySpark, Databricks, Apache Kafka, Apache Airflow, AWS (EC2, S3, VPC, IAM, RDS, Lambda, ECS, EKS)

Frameworks: Django REST Framework, Node.js/Express, React, Next.js

DevOps & Observability: Docker, Kubernetes, ELK Stack, Prometheus, Grafana, Sentry, GitHub CI/CD

Tools & Methods: Git, Power BI, PyTest, Agile/SDLC, System Design, Microservices Architecture

Certifications: AWS Solutions Architect – Associate (in progress) | Databricks Certified Data Engineer (in progress)

PROFESSIONAL EXPERIENCE

Data Engineer | Tredence Inc. | Bengaluru, India | October 2023 – Present

- Delivered clinical insights 40% faster for 200+ Banner Health stakeholders by engineering a HIPAA-compliant Medallion Architecture pipeline (Bronze → Silver → Gold) processing **10M+ patient records/day** from Cerner, Health Catalyst, McKesson, and Teradata — computing 100+ metrics (ICU/BIPAP LOS, Intubation Rate) via PySpark on Databricks with real-time Kafka streaming and Power BI reporting.
- Cut pipeline development time by 30% and eliminated manual migration risk by building **T-Ingestor**, a metadata-driven YAML-config ingestion accelerator for Airflow/Databricks Workflows; extended with an LLM-based (Claude/GPT) SQL migration pipeline using SQLGlot Transpiler and DataCompy to port Health Catalyst SQL to Spark SQL at scale.
- Enabled real-time clinical decision-making across billing, quality, and ops teams by architecting a config-driven measure engine (Base Population → Inclusion/Exclusion → Denominator → Numerator) producing standardised rate metrics by facility, provider, and cohort; led UAT with clinical SMEs before production release.
- Increased test coverage by 60% and cut QA turnaround by 35%** by building a PyTest-based automation framework covering 100+ pipeline validations with GitHub CI/CD integration across dev/QA/prod environments.

PROJECTS

Praxis — CLI SaaS Scaffolding Tool (*Top 50 Peerlist · Top 100 Product Hunt*)

- Saved **20+ dev hours per project** and accelerated SaaS prototyping by 70% by building a CLI tool that generates modular Next.js/Vite.js + Node.js apps — preconfigured with PostgreSQL/MongoDB, Prisma ORM, JWT Auth, Redis, BullMQ, Dockerised Prometheus/Grafana monitoring, and prebuilt Auth/CRUD/OAuth/LemonSqueezy routes; delivers **~1,500 QPS, <200 ms latency, and 99.95% uptime** with 80% less boilerplate. Used by 100+ developers.

Praxis Pro — Enterprise Django Boilerplate

- Reduced backend setup by 50%, developer onboarding by 35%, and manual infra work by 80% by building an enterprise-grade Django REST Framework boilerplate with JWT auth, Unicorn + Nginx containerisation, and Redoc/Swagger docs; automated Kubernetes orchestration (Pods, Ingress, Services via Kubectl/Minikube) boosting scalability by 40%, and ELK Stack + Sentry integration improving observability by 60%.

EDUCATION

Kalinga Institute of Industrial Technology, Bhubaneswar, Odisha

B.Tech – Information Technology | 8.5 CGPA | 2020 – 2024